

main-algo

March 7, 2022

```
[ ]: #importing core libraries
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import numpy as np

#importing essential libraries
from sklearn.model_selection import train_test_split
from sklearn.metrics import classification_report
from sklearn import metrics
#statistics
from scipy import stats
import statsmodels.api as sm

#importing Machine learning libraries
from sklearn.linear_model import LogisticRegression
from sklearn.svm import SVC
from sklearn.naive_bayes import GaussianNB
from sklearn.ensemble import RandomForestClassifier
from sklearn.neighbors import KNeighborsClassifier

[ ]: #importing and reading dataset

data1=pd.read_csv("student-por.csv",sep=";")
data2=pd.read_csv("student-mat.csv",sep=";")

stu=pd.concat([data1, data2])
stu["total_grades"]=(stu["G1"]+stu["G2"]+stu["G3"])/3

stu=stu.drop(["G1", "G2", "G3"],axis=1)
max=stu["total_grades"].max()
min=stu["total_grades"].min()

[ ]: #defining function for categorizing grades into 3

def marks(total_grades):
    if(total_grades<7):
```

```

        return("low")
    elif(total_grades>=7 and total_grades<14):
        return("average")
    elif(total_grades>=14):
        return("high")
stu["grades"]=stu["total_grades"].apply(marks)

```

```
[ ]: #data description
```

```

stu.dtypes
stu.describe(include="all")

```

```
[ ]:
```

	school	sex	age	address	famsize	Pstatus	Medu \
count	1044	1044	1044.000000	1044	1044	1044	1044.000000
unique	2	2	NaN	2	2	2	NaN
top	GP	F	NaN	U	GT3	T	NaN
freq	772	591	NaN	759	738	923	NaN
mean	NaN	NaN	16.726054	NaN	NaN	NaN	2.603448
std	NaN	NaN	1.239975	NaN	NaN	NaN	1.124907
min	NaN	NaN	15.000000	NaN	NaN	NaN	0.000000
25%	NaN	NaN	16.000000	NaN	NaN	NaN	2.000000
50%	NaN	NaN	17.000000	NaN	NaN	NaN	3.000000
75%	NaN	NaN	18.000000	NaN	NaN	NaN	4.000000
max	NaN	NaN	22.000000	NaN	NaN	NaN	4.000000

	Fedu	Mjob	Fjob	... romantic	famrel	freetime \
count	1044.000000	1044	1044	... 1044	1044.000000	1044.000000
unique	NaN	5	5	... 2	NaN	NaN
top	NaN	other	other	... no	NaN	NaN
freq	NaN	399	584	... 673	NaN	NaN
mean	2.387931	NaN	NaN	... NaN	3.935824	3.201149
std	1.099938	NaN	NaN	... NaN	0.933401	1.031507
min	0.000000	NaN	NaN	... NaN	1.000000	1.000000
25%	1.000000	NaN	NaN	... NaN	4.000000	3.000000
50%	2.000000	NaN	NaN	... NaN	4.000000	3.000000
75%	3.000000	NaN	NaN	... NaN	5.000000	4.000000
max	4.000000	NaN	NaN	... NaN	5.000000	5.000000

	goout	Dalc	Walc	health	absences \
count	1044.000000	1044.000000	1044.000000	1044.000000	1044.000000
unique	NaN	NaN	NaN	NaN	NaN
top	NaN	NaN	NaN	NaN	NaN
freq	NaN	NaN	NaN	NaN	NaN
mean	3.156130	1.494253	2.284483	3.543103	4.434866
std	1.152575	0.911714	1.285105	1.424703	6.210017
min	1.000000	1.000000	1.000000	1.000000	0.000000
25%	2.000000	1.000000	1.000000	3.000000	0.000000

50%	3.000000	1.000000	2.000000	4.000000	2.000000
75%	4.000000	2.000000	3.000000	5.000000	6.000000
max	5.000000	5.000000	5.000000	5.000000	75.000000

	total_grades	grades
count	1044.000000	1044
unique	NaN	3
top	NaN	average
freq	NaN	744
mean	11.267241	NaN
std	3.218805	NaN
min	1.333333	NaN
25%	9.333333	NaN
50%	11.333333	NaN
75%	13.333333	NaN
max	19.333333	NaN

[11 rows x 32 columns]

```
[ ]: #checking for info & null values
```

```
stu.info()
stu.isnull().any()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 1044 entries, 0 to 394
Data columns (total 32 columns):
#   Column          Non-Null Count  Dtype
---  -
0   school          1044 non-null   object
1   sex              1044 non-null   object
2   age              1044 non-null   int64
3   address          1044 non-null   object
4   famsize          1044 non-null   object
5   Pstatus          1044 non-null   object
6   Medu              1044 non-null   int64
7   Fedu              1044 non-null   int64
8   Mjob              1044 non-null   object
9   Fjob              1044 non-null   object
10  reason            1044 non-null   object
11  guardian          1044 non-null   object
12  traveltime        1044 non-null   int64
13  studytime          1044 non-null   int64
14  failures           1044 non-null   int64
15  schoolsup           1044 non-null   object
16  famsup             1044 non-null   object
17  paid               1044 non-null   object
```

```

18 activities      1044 non-null    object
19 nursery         1044 non-null    object
20 higher          1044 non-null    object
21 internet        1044 non-null    object
22 romantic        1044 non-null    object
23 famrel          1044 non-null    int64
24 freetime        1044 non-null    int64
25 goout           1044 non-null    int64
26 Dalc            1044 non-null    int64
27 Walc            1044 non-null    int64
28 health          1044 non-null    int64
29 absences        1044 non-null    int64
30 total_grades    1044 non-null    float64
31 grades          1044 non-null    object
dtypes: float64(1), int64(13), object(18)
memory usage: 269.2+ KB

```

```

[ ]: school      False
sex             False
age             False
address         False
famsize        False
Pstatus        False
Medu           False
Fedu           False
Mjob           False
Fjob           False
reason         False
guardian       False
traveltime     False
studytime     False
failures       False
schoolsup      False
famsup         False
paid           False
activities     False
nursery        False
higher         False
internet       False
romantic       False
famrel         False
freetime       False
goout          False
Dalc           False
Walc           False
health         False
absences       False

```

```
total_grades    False
grades          False
dtype: bool
```


```
[ ]: #visualizing the grades per number of students
```

```
plt.figure(figsize=(8,6))
sns.countplot(stu["grades"], order=["low","average","high"], palette='Set1')
plt.title('Final Grade - Number of Students',fontsize=20)
plt.xlabel('Final Grade', fontsize=16)
plt.ylabel('Number of Student', fontsize=16)
```

```
/home/el-sunais/anaconda3/lib/python3.8/site-packages/seaborn/_decorators.py:36:
FutureWarning: Pass the following variable as a keyword arg: x. From version
0.12, the only valid positional argument will be `data`, and passing other
arguments without an explicit keyword will result in an error or
misinterpretation.
```

```
warnings.warn(
```


```
[ ]: Text(0, 0.5, 'Number of Student')
```



```
[ ]: #describing by correlation matrix

corr=stu.corr()
plt.figure(figsize=(20,20))
sns.heatmap(corr, annot=True, cmap="Reds")
plt.title('Correlation Heatmap', fontsize=20)

[ ]: Text(0.5, 1.0, 'Correlation Heatmap')
```



```
[ ]: #comparing school with grades
sns.boxplot(x="school", y="total_grades", data=stu)

school_counts=stu["school"].value_counts().to_frame()
school_counts.rename(columns={"school":"school_counts"},inplace=True)
school_counts.index.name='school'


school_sns=sns.countplot(hue=stu["school"],x=stu["grades"],data=stu)


#crosstab is expanded form of value counts the the factors inside any variables
perc=(lambda col:col/col.sum())
index=["average","high","low"]
schooltab1=pd.crosstab(columns=stu.school,index=stu.grades)

school_perc=schooltab1.apply(perc).reindex(index)

school_perc.plot.bar(colormap="PiYG_r",fontsize=15,figsize=(7,7))
plt.title('Final Grade By school', fontsize=20)
plt.ylabel('Percentage of Student Counts ', fontsize=16)
plt.xlabel('Final Grade', fontsize=16)
plt.show()


#so by graph we know that school has impact on grades of students
```





```
[ ]: #comparing sex with grades
sns.boxplot(x="sex", y="total_grades", data=stu)
school_counts=stu["sex"].value_counts()


#as the graph of sex nearly overlaps so it will not have impact on grades
stu=stu.drop(["sex"],axis=1)
```





```
[ ]: #comparing address with grades
sns.boxplot(x="address", y="total_grades", data=stu)
index=["average", "high", "low"]
addresstab1=pd.crosstab(columns=stu.address, index=stu.grades)

address_perc=addresstab1.apply(perc).reindex(index)

address_perc.plot.bar(colormap="PiYG_r", fontsize=15, figsize=(7,7))
plt.title('Final Grade By address', fontsize=20)
plt.ylabel('Percentage of Student Counts ', fontsize=16)
plt.xlabel('Final Grade', fontsize=16)
plt.show()
#address is factor for the grades
```







```
[ ]: #comparing famsize with grades
sns.boxplot(x="famsize", y="total_grades", data=stu)
famsizetab1=pd.crosstab(columns=stu.famsize,index=stu.grades)

famsize_perc=famsizetab1.apply(perc).reindex(index)

famsize_perc.plot.bar(colormap="PiYG_r",fontsize=15,figsize=(7,7))
plt.title('Final Grade By famsize', fontsize=20)
plt.ylabel('Percentage of Student Counts ', fontsize=16)
```

```
plt.xlabel('Final Grade', fontsize=16)
plt.show()
#famsize has great impact on grades
```







```
[ ]: #comparing pstatus with grades
sns.boxplot(x="Pstatus", y="total_grades", data=stu)
Pstatusab1=pd.crosstab(columns=stu.Pstatus,index=stu.grades)

Pstatus_perc=Pstatusab1.apply(perc).reindex(index)

Pstatus_perc.plot.bar(colormap="PiYG_r",fontSize=15,figsize=(7,7))
plt.title('Final Grade By Pstatus', fontsize=20)
plt.ylabel('Percentage of Student Counts ', fontsize=16)
```


```
plt.xlabel('Final Grade', fontsize=16)
plt.show()
#it is not a good factor
```





```
[ ]: #comparing jobs
sns.boxplot(x="Mjob", y="total_grades", data=stu)
sns.boxplot(x="Fjob", y="total_grades", data=stu)
stu1=stu[["Fjob","Mjob","total_grades"]]
job_grp=stu1.groupby(['Mjob','Fjob'],as_index=False).mean()
job_pivot=job_grp.pivot(index='Mjob',columns='Fjob',values='total_grades')


#so father and mother jobs has great impact on grades
```



```
[ ]: #comparing reasons
sns.boxplot(x="reason", y="total_grades", data=stu)

#it has impact on the grades
```


```
[ ]: <AxesSubplot:xlabel='reason', ylabel='total_grades'>
```





```
[ ]: #comparing guardians
sns.boxplot(x="guardian", y="total_grades", data=stu)

guardiantab1=pd.crosstab(columns=stu.guardian,index=stu.grades)
guardian_perc=guardiantab1.apply(perc).reindex(index)
guardian_perc.plot.bar(colormap="BrBG",fontsize=15,figsize=(7,7))
plt.title('Final Grade By guardian', fontsize=20)
plt.ylabel('Percentage of Student Counts ', fontsize=16)
plt.xlabel('Final Grade', fontsize=16)
plt.show()

#so guardian has gr4t impact on grades
```






```
[ ]: #support of family and school
sns.boxplot(x="schoolsup", y="total_grades", data=stu)


#it is the important factor
sns.boxplot(x="famsup", y="total_grades", data=stu)
stu[["famsup", "total_grades"]].groupby(["famsup"], as_index=False).mean()


#famsup does not have great impact on grades
stu=stu.drop(["famsup"], axis=1)
```



```
[ ]: #comparing paid attributes
sns.boxplot(x="paid", y="total_grades", data=stu)
paidtab1=pd.crosstab(columns=stu.paid,index=stu.grades)
paid_perc=paidtab1.apply(perc).reindex(index)
paid_perc.plot.bar(colormap="BrBG",fontsize=15,figsize=(7,7))
plt.title('Final Grade By paid', fontsize=20)
plt.ylabel('Percentage of Student Counts ', fontsize=16)
plt.xlabel('Final Grade', fontsize=16)
plt.show()

#paid does not have much influence on grades so
stu=stu.drop(["paid"],axis=1)
```

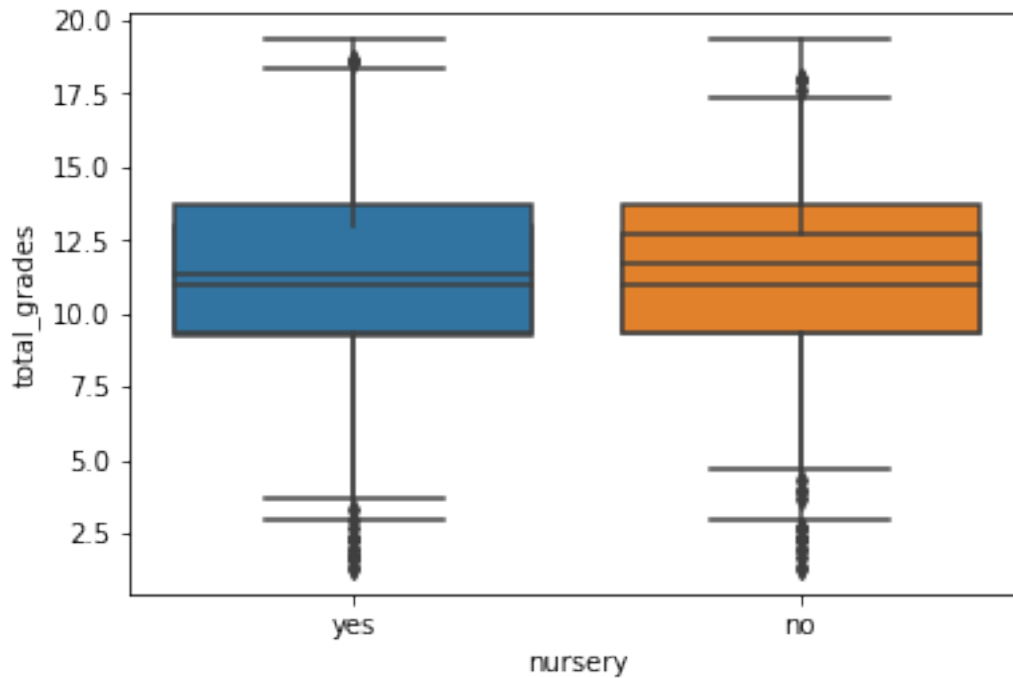




```
[ ]: sns.boxplot(x="activities", y="total_grades", data=stu)
      #is has great impact on student perforamnce

sns.boxplot(x="nursery", y="total_grades", data=stu)
      #it does not have great impact on performance

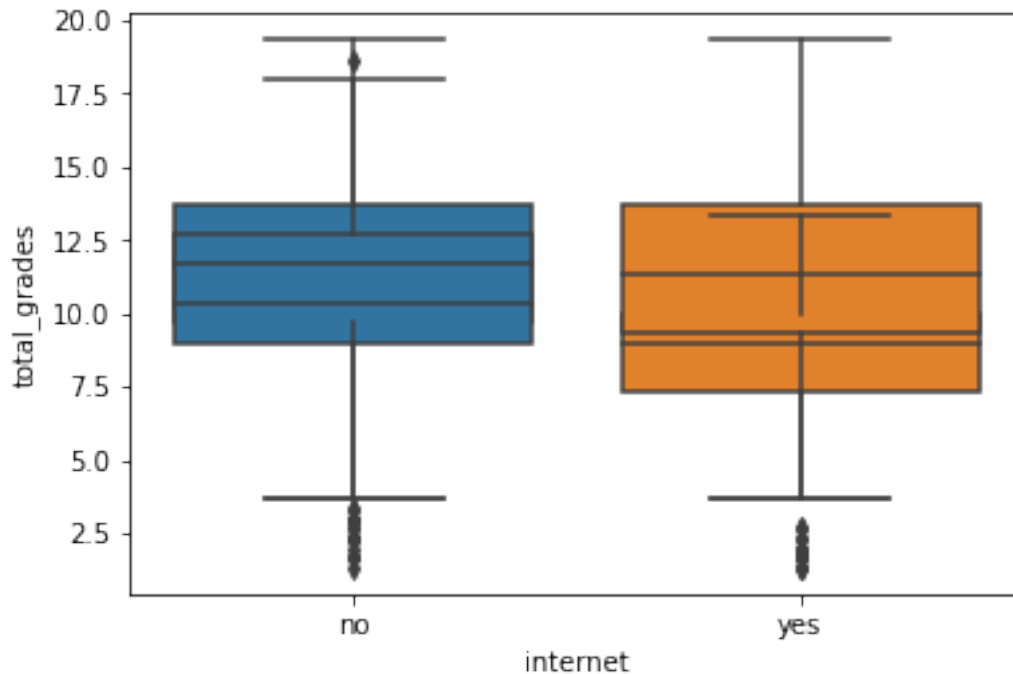
stu=stu.drop(["nursery"],axis=1)
```



```
[ ]: #comparing if higher education of students have impact on performance
sns.boxplot(x="higher", y="total_grades", data=stu)


sns.boxplot(x="internet", y="total_grades", data=stu)
#internet also have great impact on performance of individual
```


```
[ ]: <AxesSubplot:xlabel='internet', ylabel='total_grades'>
```



```
[ ]: #high school romance impact on the performance of students
sns.boxplot(x="romantic", y="total_grades", data=stu)
romantictab1=pd.crosstab(columns=stu.romantic,index=stu.grades)
romantic_perc=romantictab1.apply(perc).reindex(index)
romantic_perc.plot.bar(colormap="BrBG",fontsize=15,figsize=(7,7))
plt.title('Final Grade By romantic', fontsize=20)
plt.ylabel('Percentage of Student Counts ', fontsize=16)
plt.xlabel('Final Grade', fontsize=16)
plt.show()

#so high school romance leads to decline in performance of students
#beware of that
```



```
[ ]: #labels of remaining features
stu.columns
```

```
[ ]: Index(['school', 'age', 'address', 'famsize', 'Pstatus', 'Medu', 'Fedu',
        'Mjob', 'Fjob', 'reason', 'guardian', 'traveltime', 'studytime',
        'failures', 'schoolsup', 'activities', 'higher', 'internet', 'romantic',
        'famrel', 'freetime', 'goout', 'Dalc', 'Walc', 'health', 'absences',
        'total_grades', 'grades'],
        dtype='object')
```

```
[ ]: #creating dummies to replace string with categorical

stu1=pd.
↳get_dummies(stu,columns=["school","address","famsize","Pstatus","Mjob","Fjob","reason","gua
↳'schoolsup', 'activities', 'higher', 'internet', 'romantic' ])
test_stu1=stu1["grades"]
teststu1=stu1["total_grades"]
train_stu1=stu1.drop(['total_grades','grades'],axis=1)
train_stu=train_stu1.values
train_stu1
```

```
[ ]:      age  Medu  Fedu  traveltime  studytime  failures  famrel  freetime  \
0      18     4     4           2           2           0         4         3
1      17     1     1           1           2           0         5         3
2      15     1     1           1           2           0         4         3
3      15     4     2           1           3           0         3         2
4      16     3     3           1           2           0         4         3
..     ...     ...     ...         ...         ...         ...         ...         ...
390    20     2     2           1           2           2         5         5
391    17     3     1           2           1           0         2         4
392    21     1     1           1           1           3         5         5
393    18     3     2           3           1           0         4         4
394    19     1     1           1           1           0         3         2
```

```
      goout  Dalc  ...  schoolsup_no  schoolsup_yes  activities_no  \
0         4     1  ...             0             1             1
1         3     1  ...             1             0             1
2         2     2  ...             0             1             1
3         2     1  ...             1             0             0
4         2     1  ...             1             0             1
..     ...     ...     ...         ...         ...         ...
390     4     4  ...             1             0             1
391     5     3  ...             1             0             1
392     3     3  ...             1             0             1
393     1     3  ...             1             0             1
394     3     3  ...             1             0             1
```

```
      activities_yes  higher_no  higher_yes  internet_no  internet_yes  \
0                  0          0          1             1             0
1                  0          0          1             0             1
2                  0          0          1             0             1
3                  1          0          1             0             1
4                  0          0          1             1             0
..                 ...         ...         ...         ...         ...
390                 0          0          1             1             0
391                 0          0          1             0             1
392                 0          0          1             1             0
```

```

393          0          0          1          0          1
394          0          0          1          0          1

```

```

      romantic_no  romantic_yes
0              1              0
1              1              0
2              1              0
3              0              1
4              1              0
..          ...          ...
390            1              0
391            1              0
392            1              0
393            1              0
394            1              0

```

[1044 rows x 48 columns]

```


[ ]: #comparing age with marks
     sns.regplot(x="age",y="total_grades",data=stu)

```

```

[ ]: <AxesSubplot:xlabel='age', ylabel='total_grades'>

```



```
[ ]: #pearson coeffiecient
stu[["age","total_grades"]].corr()

#p-value
pearson_coef , p_value=stats.pearsonr(stu["age"],stu["total_grades"])
print("The Pearson Correlation Coefficient is", pearson_coef, " with a P-value_
↳ of P =", p_value)

#age is not a good factor
```

The Pearson Correlation Coefficient is -0.12913452270388814 with a P-value of P = 2.8516659922093794e-05

```
[ ]: #using backward elimination for finding optimal featrures

#if p-value is greater than 0.6 than we will removethat feature
X=np.append(arr=np.ones((1044,1)).astype(int),values=train_stu,axis=1)
X_opt = X[:, [0, 1, 2, 3, 4,5,6,7,8,9,10,11,12,13]]
regressor_ols=sm.OLS(endog=teststu1,exog=X_opt).fit()
regressor_ols.summary()

X_opt = X[:, [0,2,3,4,5,6,7,8,9,10,11,12,13]]
regressor_ols=sm.OLS(endog=teststu1,exog=X_opt).fit()
regressor_ols.summary()

X_opt = X[:, [0,2,3,4,5,6,7,9,10,11,12,13]]
regressor_ols=sm.OLS(endog=teststu1,exog=X_opt).fit()
regressor_ols.summary()
```

```
[ ]: <class 'statsmodels.iolib.summary.Summary'>
"""
```

```

                                OLS Regression Results
=====
Dep. Variable:                total_grades    R-squared:                0.215
Model:                            OLS        Adj. R-squared:            0.207
Method:                    Least Squares    F-statistic:                25.74
Date:                Mon, 07 Mar 2022    Prob (F-statistic):        1.53e-47
Time:                        16:50:04    Log-Likelihood:            -2574.8
No. Observations:                1044    AIC:                        5174.
Df Residuals:                    1032    BIC:                        5233.
Df Model:                        11
Covariance Type:                nonrobust
=====

```

	coef	std err	t	P> t	[0.025	0.975]
const	10.9881	0.646	17.005	0.000	9.720	12.256
x1	0.3571	0.105	3.394	0.001	0.151	0.564

x2	0.1062	0.106	0.998	0.318	-0.103	0.315
x3	-0.1982	0.126	-1.570	0.117	-0.446	0.050
x4	0.4150	0.111	3.743	0.000	0.197	0.632
x5	-1.6060	0.142	-11.344	0.000	-1.884	-1.328
x6	0.1104	0.097	1.133	0.258	-0.081	0.302
x7	-0.1906	0.085	-2.235	0.026	-0.358	-0.023
x8	-0.2267	0.126	-1.798	0.072	-0.474	0.021
x9	0.0377	0.096	0.391	0.696	-0.152	0.227
x10	-0.1392	0.064	-2.190	0.029	-0.264	-0.014
x11	-0.0185	0.015	-1.269	0.205	-0.047	0.010

```
=====
Omnibus:                40.630    Durbin-Watson:                1.886
Prob(Omnibus):           0.000    Jarque-Bera (JB):         48.582
Skew:                    -0.426    Prob(JB):                 2.82e-11
Kurtosis:                3.625    Cond. No.                 73.3
=====
```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

"""

```
[ ]: #now we merge our training data
train_x=np.concatenate((X_opt,X[:,14:49]),axis=1)
stu[["Medu","total_grades"]].corr()
stu[["Fedu","total_grades"]].corr()
```

```
[ ]:          Fedu  total_grades
Fedu          1.00000      0.18661
total_grades  0.18661      1.00000
```

```
[ ]: #final feature extraction cleaning
train_stu2=train_stu1.drop(["age","freetime"],axis=1)
np1=[1 for i in range(0,1044)]
train_stu2.insert(loc=0,column= "noimprotance", value=np1)
```

```
[ ]: #now after getting the proper features we will split the data
train_stu2.columns
train_stu2.head(30)
```

```
[ ]:   noimprotance  Medu  Fedu  traveltime  studytime  failures  famrel  goout  \
0                1    4    4             2           2           0     4     4
1                1    1    1             1           2           0     5     3
2                1    1    1             1           2           0     4     2
3                1    4    2             1           3           0     3     2
4                1    3    3             1           2           0     4     2
5                1    4    3             1           2           0     5     2
```

6	1	2	2	1	2	0	4	4
7	1	4	4	2	2	0	4	4
8	1	3	2	1	2	0	4	2
9	1	3	4	1	2	0	5	1
10	1	4	4	1	2	0	3	3
11	1	2	1	3	3	0	5	2
12	1	4	4	1	1	0	4	3
13	1	4	3	2	2	0	5	3
14	1	2	2	1	3	0	4	2
15	1	4	4	1	1	0	4	4
16	1	4	4	1	3	0	3	3
17	1	3	3	3	2	0	5	2
18	1	3	2	1	1	3	5	5
19	1	4	3	1	1	0	3	3
20	1	4	3	1	2	0	4	1
21	1	4	4	1	1	0	5	2
22	1	4	2	1	2	0	4	1
23	1	2	2	2	2	0	5	4
24	1	2	4	1	3	0	4	2
25	1	2	2	1	1	0	1	2
26	1	2	2	1	1	0	4	2
27	1	4	2	1	1	0	2	4
28	1	3	4	1	2	0	5	3
29	1	4	4	1	2	0	4	5

	Dalc	Walc	...	schoolsup_no	schoolsup_yes	activities_no	\
0	1	1	...	0	1	1	
1	1	1	...	1	0	1	
2	2	3	...	0	1	1	
3	1	1	...	1	0	0	
4	1	2	...	1	0	1	
5	1	2	...	1	0	0	
6	1	1	...	1	0	1	
7	1	1	...	0	1	1	
8	1	1	...	1	0	1	
9	1	1	...	1	0	0	
10	1	2	...	1	0	1	
11	1	1	...	1	0	0	
12	1	3	...	1	0	0	
13	1	2	...	1	0	1	
14	1	1	...	1	0	1	
15	1	2	...	1	0	1	
16	1	2	...	1	0	0	
17	1	1	...	0	1	0	
18	2	4	...	1	0	0	
19	1	3	...	1	0	0	
20	1	1	...	1	0	1	

21	1	1	...	1	0	1
22	1	3	...	1	0	0
23	2	4	...	1	0	0
24	1	1	...	0	1	0
25	1	3	...	1	0	1
26	1	2	...	1	0	1
27	2	4	...	1	0	1
28	1	1	...	0	1	0
29	5	5	...	1	0	0

	activities_yes	higher_no	higher_yes	internet_no	internet_yes	\
0	0	0	1	1	0	
1	0	0	1	0	1	
2	0	0	1	0	1	
3	1	0	1	0	1	
4	0	0	1	1	0	
5	1	0	1	0	1	
6	0	0	1	0	1	
7	0	0	1	1	0	
8	0	0	1	0	1	
9	1	0	1	0	1	
10	0	0	1	0	1	
11	1	0	1	0	1	
12	1	0	1	0	1	
13	0	0	1	0	1	
14	0	0	1	0	1	
15	0	0	1	0	1	
16	1	0	1	0	1	
17	1	0	1	1	0	
18	1	0	1	0	1	
19	1	0	1	0	1	
20	0	0	1	0	1	
21	0	0	1	0	1	
22	1	0	1	0	1	
23	1	0	1	0	1	
24	1	0	1	0	1	
25	0	0	1	0	1	
26	0	0	1	0	1	
27	0	0	1	0	1	
28	1	0	1	0	1	
29	1	0	1	0	1	

	romantic_no	romantic_yes
0	1	0
1	1	0
2	1	0
3	0	1

4	1	0
5	1	0
6	1	0
7	1	0
8	1	0
9	1	0
10	1	0
11	1	0
12	1	0
13	1	0
14	0	1
15	1	0
16	1	0
17	1	0
18	1	0
19	1	0
20	1	0
21	1	0
22	1	0
23	1	0
24	1	0
25	1	0
26	1	0
27	1	0
28	1	0
29	0	1

[30 rows x 47 columns]

```
[ ]: #splitting dataset into training and testing data
```

```
X_train, X_test, y_train, y_test = train_test_split(train_x, test_stu1,
    ↪ test_size = 0.2, random_state = 0)
```

```
[ ]: #features by importance using RFC
```

```
rf_classifier = RandomForestClassifier(n_estimators=80, criterion="entropy", random_state=0)
```

```
rf_classifier.fit(X_train, y_train)
```

```
sv_classifier = SVC(kernel = 'rbf', random_state = 0)
```

```
sv_classifier.fit(X_train, y_train)
```

```
lr_classifier = LogisticRegression(random_state=0)
```

```
lr_classifier.fit(X_train, y_train)
```

```
kn_classifier = KNeighborsClassifier(n_neighbors=5, metric='minkowski', p=2)
```

```
kn_classifier.fit(X_train, y_train)
```

```
nb_classifier = GaussianNB()
```

```
nb_classifier.fit(X_train, y_train)
```

```
/home/el-sunais/anaconda3/lib/python3.8/site-  
packages/sklearn/linear_model/_logistic.py:763: ConvergenceWarning: lbfgs failed  
to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
```

```
[ ]: GaussianNB()
```

```
[ ]: #Predictions Performance Scores
```

```
rf_score = rf_classifier.score(X_test, y_test)  
sv_score = sv_classifier.score(X_test, y_test)  
lr_score = lr_classifier.score(X_test, y_test)  
kn_score = kn_classifier.score(X_test, y_test)  
nb_score = nb_classifier.score(X_test, y_test)  
  
print('-----Prediction Scores-----')  
  
print('Random Forest Prediction Score is', rf_score*100,'%')  
print('Support Vector Prediction Score is', sv_score*100,'%')  
print('Logistic Regression Prediction Score is', lr_score*100,'%')  
print('K-Nearest Classifier Prediction Score is', kn_score*100,'%')  
print('Naive Bayes Prediction Score is', nb_score*100,'%')
```

```
-----Prediction Scores-----
```

```
Random Forest Prediction Score is 76.55502392344498 %  
Support Vector Prediction Score is 73.20574162679426 %  
Logistic Regression Prediction Score is 75.5980861244019 %  
K-Nearest Classifier Prediction Score is 71.29186602870813 %  
Naive Bayes Prediction Score is 33.49282296650718 %
```

```
[ ]: #predicting the test set results
```

```
yrf = rf_classifier.predict(X_test)  
ysv = sv_classifier.predict(X_test)  
ylr = lr_classifier.predict(X_test)  
ykn = kn_classifier.predict(X_test)  
ynb = nb_classifier.predict(X_test)
```

```
#determining the precision, recall and f1-score
```

```

RF_report=classification_report(y_test,yrf)
SV_report=classification_report(y_test,ysv)
LR_report=classification_report(y_test,ylr)
KN_report=classification_report(y_test,ykn)
NB_report=classification_report(y_test,ynb)

```

```

/home/el-sunais/anaconda3/lib/python3.8/site-
packages/sklearn/metrics/_classification.py:1245: UndefinedMetricWarning:
Precision and F-score are ill-defined and being set to 0.0 in labels with no
predicted samples. Use `zero_division` parameter to control this behavior.
  _warn_prf(average, modifier, msg_start, len(result))
/home/el-sunais/anaconda3/lib/python3.8/site-
packages/sklearn/metrics/_classification.py:1245: UndefinedMetricWarning:
Precision and F-score are ill-defined and being set to 0.0 in labels with no
predicted samples. Use `zero_division` parameter to control this behavior.
  _warn_prf(average, modifier, msg_start, len(result))
/home/el-sunais/anaconda3/lib/python3.8/site-
packages/sklearn/metrics/_classification.py:1245: UndefinedMetricWarning:
Precision and F-score are ill-defined and being set to 0.0 in labels with no
predicted samples. Use `zero_division` parameter to control this behavior.
  _warn_prf(average, modifier, msg_start, len(result))

```

```

[ ]: print('Random Forest Classifier Report:')
     print(RF_report)

```

```

Random Forest Classifier Report:

```

	precision	recall	f1-score	support
average	0.79	0.93	0.85	153
high	0.67	0.42	0.52	38
low	0.40	0.11	0.17	18
accuracy			0.77	209
macro avg	0.62	0.49	0.51	209
weighted avg	0.73	0.77	0.73	209

```

[ ]: print('Support Vector Classifier Report:')
     print(SV_report)

```

```

Support Vector Classifier Report:

```

	precision	recall	f1-score	support
average	0.73	1.00	0.85	153
high	0.00	0.00	0.00	38
low	0.00	0.00	0.00	18
accuracy			0.73	209

macro avg	0.24	0.33	0.28	209
weighted avg	0.54	0.73	0.62	209

```
[ ]: print('Logistic Regression Classifier Report:')
      print(LR_report)
```

Logistic Regression Classifier Report:

	precision	recall	f1-score	support
average	0.77	0.95	0.85	153
high	0.69	0.29	0.41	38
low	0.40	0.11	0.17	18
accuracy			0.76	209
macro avg	0.62	0.45	0.48	209
weighted avg	0.72	0.76	0.71	209

```
[ ]: print('K-Nearest Classifier Report:')
      print(KN_report)
```

K-Nearest Classifier Report:

	precision	recall	f1-score	support
average	0.76	0.90	0.83	153
high	0.41	0.29	0.34	38
low	0.00	0.00	0.00	18
accuracy			0.71	209
macro avg	0.39	0.40	0.39	209
weighted avg	0.63	0.71	0.67	209

```
[ ]: print('Naive Bayes Classifier Report:')
      print(NB_report)
```

Naive Bayes Classifier Report:

	precision	recall	f1-score	support
average	0.85	0.18	0.30	153
high	0.24	0.97	0.39	38
low	0.22	0.28	0.24	18
accuracy			0.33	209
macro avg	0.44	0.48	0.31	209
weighted avg	0.68	0.33	0.31	209